

Each anime is modelled as its own row in a single Anime table, utilising an auto incrementing anime_id SERIAL primary key. This surrogate key ensures a unique, stable identifier, for example even when two shows share the same name. All single-valued properties - english_name, score, type, episodes, aired, premiered, source, duration, rating, ranked and popularity live as columns on Anime table.

As every show may have many genres, studios and producers - and each genre, studio or producer may apply to many shows, we pulled those lists out into three lookup tables: Genre, Studio and Producer. Every lookup table has its own surrogate primary key and a name column to ensure we store 'Action' for example, only once and refrain from repeating text across multiple anime rows.

In order to connect Anime to Genre, Anime to Studio, and Anime to producer in many-to-many format, three join tables are created. These three join tables Anime_Genre, Anime_Studio, Anime_Producer each with its own primary key and two foreign keys back to the parent tables. ON DELETE CASCADE is included in order to delete an anime, automatically cleaning up its related join entries.

